AI VOICE AGENT PLATFORM

Client Proposal — Service Plans & Pricing

Prepared for

EMPOWER AI 365

AI Voice Agent Platform | 2026

Confidential — For Discussion Only

01 | WHAT IS THIS PRODUCT?

Think of this as your own AI receptionist business — fully built, fully branded under your name, and ready to sell to any client who answers phones.

Your clients — whether they run a dental clinic, a real estate agency, a law firm, or a call centre — subscribe to your service. From day one, an AI agent answers every call for them, 24 hours a day, 7 days a week. No human on the line. No missed calls. No voicemail black holes.

You set the price. You hold the client relationship. You keep the margin. Techvera builds and runs the technology behind it — completely invisible to your clients.

> You are not reselling someone else's product. You own this platform. Your brand is on it. Your clients never see Techvera or any third-party name — ever.

02 | WHAT DOES THE AI ACTUALLY DO FOR YOUR CLIENTS?

When your client's phone rings — whether it is 9am or 3am — the AI answers. It speaks naturally, handles the conversation, and takes action. Here is what it does on every call:

📞 Answers every call — instantly, 24/7

No hold music. No voicemail. The AI greets the caller, introduces the business, and gets straight to helping them.

📅 Books appointments in real time

The AI checks your client's live calendar and books, reschedules, or cancels appointments during the call — sending an instant SMS or email confirmation to the caller.

💬 Answers questions from a knowledge base

You upload your client's FAQs, services, and documents. The AI answers from that content only — it never guesses. If it does not know the answer, it escalates to a real person.

🎯 Identifies and logs new leads

When a caller is a potential new customer, the AI detects it, scores the lead, and instantly notifies your client's sales team via Slack, email, or SMS — with a full call summary attached.

🔀 Routes calls to the right department

The AI understands why someone is calling and connects them to the right person or team — automatically. Warm transfer (agent introduces the caller) or cold transfer (instant handoff).

🚨 Escalates emergencies immediately

If a caller uses an emergency phrase, the AI contacts your client's on-call person directly. Multiple contacts are tried in sequence until someone responds.

📊 Records, transcribes, and summarises every call

Every call is recorded and stored. A word-for-word transcript plus an AI summary — caller intent, outcome, action items, sentiment score, and metadata — delivered to Slack, email, or SMS.

📤 Dials outbound contacts automatically

The AI dials from a contact list, detects voicemail, leaves a message or schedules a callback — no human dialler needed.

03 | THE THREE PLANS — WHICH ONE IS RIGHT FOR YOU?

There are three ways to own this platform. Each gives you the same AI capabilities — the difference is in how the technology is set up behind the scenes, how quickly you can start selling, and how much it costs per minute at scale.

Tier 1	Tier 2	Tier 3
Quick Start	**Full Control**	**Best of Both**
"I want to start selling this week"	*"I want the lowest possible cost at scale and own my margins"*	*"I want to serve every type of client that gives flexibility of both TIER 1 and TIER 2 under one platform"*
$1,500	**$2,500**	**$3,000**
licence	licence	licence

04 | TIER 1 — QUICK START | $1,500 Licence

In plain English — what is Tier 1?

Tier 1 is the fastest way to go from zero to a live, revenue-generating product. Techvera handles every technical aspect of running the platform — you focus entirely on selling it.

Your clients sign up, plug in their phone number, and the AI is live within days. The technology runs on a professional managed service called Retell AI — the same infrastructure used by enterprise companies globally. You never touch a server. You never manage any software. If something needs fixing, Techvera fixes it.

The trade-off: because someone else is managing the infrastructure for you, the per-minute cost is slightly higher than Tier 2. But margins are still strong — and for most new businesses, the speed and simplicity of Tier 1 is exactly the right starting point.

Who should choose Tier 1?

- You want to start selling within the next 1–2 weeks
- You do not have a technical team — and you do not want one
- You are signing your first clients and want a reliable, proven setup
- Your clients are small to mid-sized businesses — clinics, agencies, service businesses
- You want Techvera to handle uptime, maintenance, and platform issues

What you get

What	Detail
Your brand	Your clients see only your brand name — no Techvera, no third-party names
Live in	3 days for a simple setup. 5–7 days for custom call routing
AI response speed	0.9–1.6 seconds average — callers do not notice the difference from a human

Voice quality	ElevenLabs — the highest-quality AI voice available today
Uptime	99.9% target — fewer than 9 hours of unplanned downtime per year
Languages	Full quality: English, Arabic (Gulf/MSA), Spanish, French. 32 more available
Calendar booking	Google Calendar, Calendly — live booking during the call, no double-booking
CRM connections	HubSpot, GoHighLevel — leads detected in real time and pushed automatically
Call recording	Every call is recorded and stored for quality review
Transcripts & summaries	Word-for-word transcript + AI summary — intent, outcome, action items, sentiment score — delivered to Slack, email, or SMS
Outbound dialler	Dials from a contact list automatically, detects voicemail, leaves message or schedules callback
Call routing	Intent-based routing — warm transfer (agent introduces caller) or cold transfer (instant handoff). No-answer fallback routes to voicemail or offers callback
Concurrent calls	20 simultaneous calls included. Additional at $8 per line per month
Who manages the tech	Techvera — fully managed for you
Compliance	SOC2 Type 1&2 and HIPAA-ready — suitable for healthcare and finance clients

What does it cost you per minute?

Your cost is made up of several components — Techvera bundles these so you pay one rate, but here is exactly what goes into it:

Component	Cost / min	What it does
Platform + speech recognition	$0.07–0.08	Retell AI orchestration, turn-taking, interruption handling. Deepgram speech-to-text bundled in.
AI brain — budget	$0.003–0.006	GPT-5 nano / Gemini Flash. Best for FAQ, scheduling, simple routing.
AI brain — recommended	$0.020–0.080	GPT-4o / Claude Sonnet. For lead qualification, complex conversations, multi-step flows.
Voice — standard	$0.015	Cartesia / Retell native voices. Good quality, lower cost.
Voice — premium	$0.040–0.070	ElevenLabs. Highest natural quality — recommended for all client-facing deployments.
Telephony (phone lines)	$0.010–0.020	Twilio / Telnyx inbound + outbound routing. Gulf/Pakistan: $0.008–0.015/min — verify before pricing.

Automation (CRM + alerts)	$0.002–0.005	Make.com / n8n webhooks, CRM sync, notifications.

Your all-in total by quality level:

Quality Level	Your Cost / min	Sell at	Your Margin	Best used when
Standard (good quality)	~$0.10–0.11	$0.18–0.20	39–45%	High volume, simple scripts
Recommended (great quality)	~$0.13–0.17	$0.22–0.25	24–40%	Most client deployments
Premium (best quality)	~$0.19–0.22	$0.24–0.28	20–28%	Complex sales, healthcare

Always price based on the higher cost estimate, not the lower. If your actual cost comes in lower, that is extra profit — not a number to plan around.

Monthly subscription plans (what you charge your clients)

Plan	Minutes / Month	Price per Minute	Monthly Total	Extra Minutes
Starter	500	$0.22	$110 / month	$0.24 / min
Growth	2,000	$0.20	$400 / month	$0.22 / min
Business	5,000	$0.18	$900 / month	$0.20 / min
Scale	15,000	$0.16	$2,400 / month	$0.18 / min

Unused minutes do not roll over. All plans include unlimited inbound and outbound calls within the allocated minutes.

05 | TIER 2 — FULL CONTROL | $2,500 Licence

In plain English — what is Tier 2?

Tier 2 is for when you are ready to own the technology outright — not just the brand, but the actual servers, the software, and every component underneath. Instead of paying a managed service like Retell AI a platform fee on every minute, you run the equivalent technology yourself on your own servers.

The result is a significantly lower cost per minute — roughly half the cost of Tier 1 at scale. You also get a multi-tenant platform, meaning one installation of the software can serve unlimited clients, all under your brand, all from a single management dashboard.

The trade-off: Tier 2 requires a developer to maintain the infrastructure. Techvera builds and configures everything — but ongoing maintenance requires a technical resource on your side or a maintenance agreement with Techvera. Setup also takes 2–4 weeks instead of days.

Who should choose Tier 2?

- You have high-volume clients — thousands of minutes per month
- You want the lowest possible per-minute cost to maximise margin
- You want to serve unlimited clients from a single platform
- You have access to a developer (in-house or via Techvera's maintenance plan)
- You are building a long-term, scalable AI services business
- Some of your clients want to know they own their own infrastructure

What you get

What	Detail
Your brand	Full white-label — zero vendor visibility anywhere in the platform
Live in	14–18 days. Add 5–7 days buffer for infrastructure setup

AI response speed	0.7–1.2 seconds — slightly faster than Tier 1 in a well-tuned deployment
Voice quality	Cartesia Sonic — very good quality, low latency (~75ms), optimised for phone calls
Uptime	99.9% target — you own the infrastructure; Techvera supports it
Languages	Full quality: English, Arabic (Gulf/MSA), Spanish, French. 32 more available
Calendar booking	Google Calendar, Outlook, and any other calendar system
CRM connections	HubSpot, GoHighLevel, Salesforce, or any custom CRM
Call recording	Every call is recorded and stored for quality review
Transcripts & summaries	Word-for-word transcript + AI summary — intent, outcome, action items, sentiment score — delivered to Slack, email, or SMS
Outbound dialler	Dials from a contact list automatically, detects voicemail, leaves message or schedules callback
Call routing	Intent-based routing — warm transfer (agent introduces caller) or cold transfer (instant handoff). No-answer fallback routes to voicemail or offers callback
Concurrent calls	Scales with your server capacity — no per-line fees
Number of clients	Unlimited — all managed from one platform dashboard
AI model choice	Access to 500+ AI models — match the model to each client's needs
Who manages the tech	You (developer required) — Techvera maintenance plan available
Compliance	Depends on your server/hosting configuration

What does it cost you per minute?

Because you deal directly with each component provider — no middleman platform fee — your per-minute cost is significantly lower. Here is the full breakdown:

Component	Cost / min	What it does
Pipeline orchestration	Free	Pipecat (open-source) — handles turn-taking, interruptions, voice activity detection. No licence fee; requires a developer to maintain.
Audio transport	$0.002–0.005	LiveKit (self-hosted) — WebRTC + SIP audio routing. Self-hosted = lowest cost.
AI brain (LLM)	$0.003–0.020	OpenRouter — budget: Gemini Flash ~$0.003/min. Recommended: GPT-4o ~$0.012–0.020/min.

Voice synthesis	$0.010–0.015	Cartesia Sonic — low latency, natural voice, phone-optimised.
Speech recognition	$0.005–0.008	Deepgram Nova-3 direct API — ~6.8% word error rate. Strongest on EN, AR, ES, FR.
Telephony (phone lines)	$0.005–0.015	Telnyx SIP. US/UK: ~$0.005/min. Gulf/South Asia/Pakistan: $0.008–0.015/min — always verify regional rates.
Automation (CRM + alerts)	$0.002–0.005	n8n / Make.com webhooks, CRM sync, notifications.
Server hosting	$0.002–0.004	AWS / DigitalOcean VPS — handles ~15–25 concurrent calls per server. Scale by adding servers as volume grows.

Your all-in total by configuration:

Configuration	Your Cost / min	Sell at	Your Margin	Best used when
Budget config	~$0.027–0.040	$0.07–0.09	38–63%	Simple scripts, very high volume
Recommended config	~$0.038–0.055	$0.09–0.12	38–56%	Most business deployments

Always price using $0.055/min as your cost ceiling. Actual costs in most deployments come in lower — that difference is your bonus margin.

Monthly subscription plans (what you charge your clients)

Plan	Minutes / Month	Price per Minute	Monthly Total	Extra Minutes
Starter	1,000	$0.09	$90 / month	$0.10 / min
Growth	5,000	$0.08	$400 / month	$0.09 / min
Business	15,000	$0.07	$1,050 / month	$0.08 / min
Scale	50,000	$0.06	$3,000 / month	$0.07 / min
Enterprise	100,000+	Custom	Custom	Custom

Unused minutes do not roll over. All plans include unlimited inbound and outbound calls within the allocated minutes.

06 | TIER 3 — BEST OF BOTH | $3,000 Licence

In plain English — what is Tier 3?

Tier 3 gives you both platforms — Tier 1 and Tier 2 — running together under one roof. From your dashboard, you manage all your clients in one place. For each client, you choose which infrastructure they run on based on what makes the most sense for them.

A new client who needs to go live this week and has strict compliance requirements? They go on the Tier 1 (Retell AI managed) layer. A high-volume enterprise client who wants the lowest cost per minute and does not need SOC2? They go on the Tier 2 (self-hosted) layer. Both clients are on your platform, under your brand.

This is the most flexible option — and the most powerful. It allows you to sell to any type of client at any scale, without ever outgrowing your own platform.

Who should choose Tier 3?

- You plan to serve a mixed portfolio of clients — small businesses and large enterprises
- You want to start selling quickly (via Tier 1) while building out the lower-cost layer (Tier 2)
- You want a single platform to manage all clients — regardless of their size or infrastructure
- You want to grow a client from small to large without migrating them off your platform
- You want maximum flexibility when pricing different clients differently

How does Tier 3 decide which layer a client goes on?

You decide — at the point of onboarding each new client — using this simple guide:

Factor	Route to Tier 1 (Managed) if...	Route to Tier 2 (Self-Hosted) if...
Speed	They need to go live within the week	They can wait 2–4 weeks for setup
Volume	Under ~10,000 minutes per month	High volume — cost per minute matters

Compliance	They need SOC2 / HIPAA compliance	Standard business — no special compliance required
Ownership	They are happy on a managed service	They want full ownership of their infrastructure
Voice	Highest quality voice is important	Good quality voice is sufficient

Pricing in Tier 3

Each client is priced individually based on which layer they are on. There is no extra platform fee for Tier 3 — the $3,000 licence covers the full build of both layers and the unified management platform.

Layer	Your Cost / min	Sell at	Margin	Use this layer for
Managed (Retell AI)	~$0.13–0.17	$0.22–0.25	24–40%	Fast deploys, SOC2, small–mid volume
Self-Hosted	~$0.038–0.055	$0.09–0.12	38–56%	High volume, cost-sensitive clients

The key advantages of Tier 3

- One dashboard. Every client. No matter how many you have or which layer they are on.
- Start small and scale. A client on Tier 1 today can move to Tier 2 when their volume grows — without leaving your platform or changing their experience.
- Sell to anyone. From a one-location dental clinic to a 50-seat outbound sales team — the platform handles both.
- One brand. Across everything. Your clients always see your name.

07 | SIDE-BY-SIDE COMPARISON

Not sure which tier is right for you? This table covers every key difference.

	Tier 1 — Quick Start	Tier 2 — Full Control	Tier 3 — Best of Both
Licence fee	$1,500	$2,500	$3,000
Ready to sell in	3–7 days	14–25 days	14–25 days
Cost per minute (typical)	$0.13–0.17	$0.038–0.055	Depends on client tier
Your sell price (suggested)	$0.22–0.25 / min	$0.09–0.12 / min	Set per client
Your margin (realistic)	24–40%	38–56%	Optimised per client
Number of clients	Separate setup per client	Unlimited on one platform	Unlimited on one platform
Client branding	Your brand only	Your brand only	Your brand only
Voice quality	Highest (ElevenLabs)	Very good (Cartesia)	Both available
Compliance (SOC2/HIPAA)	Yes — built in	Depends on your setup	Yes — on managed clients
Best for	Selling your first clients fast	High-volume / enterprise clients	Selling to all client types

08 | IMPORTANT THINGS TO KNOW

Arabic language support

The AI speaks English, Spanish, and French at full production quality. For Arabic — Gulf Arabic and Modern Standard Arabic are fully supported. If your clients serve callers from Lebanon, Syria, Jordan, or North Africa, accent-specific testing is required before going live in those regions. This is a one-time step Techvera handles at setup.

CRM connections

HubSpot, GoHighLevel, and Salesforce are tested and ready to connect. If your client uses a different CRM system, Techvera will scope and build that connection — quoted separately. CRM connections are live integrations, not one-time setups. They need a maintenance plan to stay working as CRM systems update.

Emergency escalation

The AI can detect emergency phrases and immediately call or text your client's on-call person. This works reliably for standard business deployments. If a client is in a life-safety-critical industry (e.g. mental health, emergency medical services), Techvera recommends a secondary human failsafe alongside the AI escalation.

Concurrent calls

On Tier 1, your account includes 20 simultaneous calls at no extra cost. If a client expects more than 20 calls at the same time, additional lines are available at $8 per line per month. On Tier 2 and Tier 3, capacity scales automatically with your server — there is no per-line fee.

Uptime

All tiers target 99.9% uptime — fewer than 9 hours of unplanned downtime per year. On Tier 1, Retell AI's enterprise infrastructure backs this commitment. On Tier 2 and Tier 3, Techvera configures this to the same standard at build time.

09 | HOW TO GET STARTED

From the moment you confirm your tier, here is what happens:

1. **Select your tier**
   Choose Tier 1, Tier 2, or Tier 3 based on where you are today and where you want to go. Not sure? Get on a call with Techvera — this decision can be made in 30 minutes.
2. **Onboarding call (30 min)**
   Techvera walks through your industry, your first clients, the CRM and calendar systems involved, and your go-live timeline. This call sets everything in motion.
3. **Agent training**
   You provide the content — your services list, FAQs, tone of voice, scripts. Techvera configures the AI agent, builds its knowledge base, and sets the guardrails so it never goes off-brand.
4. **Connections built and tested**
   Calendar, CRM, Slack/email alerts, emergency escalation, and call routing are all connected and fully tested before any real caller hears the AI.
5. **Quality review**
   Every feature is tested on a live phone number — including edge cases. Nothing goes live until it has passed a structured sign-off.
6. **Your product is live**
   The AI is live on your client's phone number, handling real calls. Techvera monitors the platform and a maintenance plan is agreed before launch.

Ready to build your AI voice business?

Every tier comes with full support through setup, testing, and go-live. You focus on sales. Techvera handles the rest.

techveraglobal@gmail.com | www.techveraglobal.com